

Dongqi Cai (蔡栋琪)

PhD Student (Fourth Year)

School of Computer Science, Beijing University of Posts and Telecommunications, China

Email: cdq@bupt.edu.cn

Homepage: <http://www.caidongqi.com/>

Research Interests

Federated Learning, Efficient NLP System, Speech Privacy.

Education

09/2024 – present	Visiting PhD , University of Cambridge <ul style="list-style-type: none">• Advisor: Nicholas D. Lane
09/2021 – present	PhD in Computer Science and Technology , BUPT <ul style="list-style-type: none">• Advisor: Shangguang Wang, Mengwei Xu• Remote Advisor: Felix Xiaozhu Lin (University of Virginia)
09/2017 – 07/2021	BS in Communication Engineering , BUPT <ul style="list-style-type: none">• Advisor: Lin Fan

Internship

07/2021 – 12/2021	Research Intern , WeBank <ul style="list-style-type: none">• Mentor: Lixin Fan
-------------------	---

Honors & Awards

- Young Elite Scientists Sponsorship (PhD student Special Program), CAST, 2025
- National Scholarship, Ministry of Education, 2024
- St John's College Fellow-Sponsored Member, University of Cambridge, 2024
- Scholar Award, NeurIPS, 2024
- CSC Scholarship, China Scholarship Council, 2024
- Travel Grant, EuroSys/MobiSys/ATC, 2024
- National Scholarship, Ministry of Education, 2023
- Outstanding Graduate Student, BUPT, 2023
- Excellent Ph.D. Students Foundation, BUPT, 2023
- Outstanding Graduate Student, State Key Laboratory of Networking and Switching Technology, 2022/2023

Academic Services

- **TPC Member**

MobiSys'24 AE, MobiCom'24 AE, NCSC-edge'22, TURC-SIGBED-China'23

- **Reviewer**

Scientific Reports, TSC, TMC, TKDE, TECS, IoTJ, SAGC'22, ICASSP'24, ICASSP'25.

- **External Reviewer**

MLSys'25, ICWS'24, IEEE EDGE'24, IEEE EDGE'23, ICWS'23, EIS'21

Conference Publications (* = equal contributions; # = corresponding)

(full list at <https://scholar.google.com/citations?user=dlimkboAAAAJ&hl=zh-CN>)

[C10] "DEPT: Decoupled Embeddings for Pre-training Language Models"

Alex Iacob, Lorenzo Sani, Meghdad Kurmanji, William F. Shen, Xinchu Qiu, **Dongqi Cai**, Yan Gao, Nicholas Donald Lane, in the *Thirteenth International Conference on Learning Representations (ICLR)*, 2025.

[C9] "ShortcutsBench: A Large-Scale Real-world Benchmark for API-based Agents"

Haiyang Shen, Yue Li, Desong Meng, **Dongqi Cai**, Sheng Qi, Li Zhang, Mengwei Xu, Yun Ma, in the *Thirteenth International Conference on Learning Representations (ICLR)*, 2025.

[C8] "SILENCE: Protecting privacy in offloaded speech understanding on wimpy devices"

Dongqi Cai, Shangguang Wang, Zeling Zhang, Felix Xiaozhu Lin, Mengwei Xu, in the *Annual Conference on Neural Information Processing Systems (NeurIPS, CCF-A)*, 2024.

[C7] "FwdLLM: Efficient Federated Finetuning of Large Language Models with Perturbed Inferences"

Mengwei Xu (My advisor), **Dongqi Cai**[#], Yaozong Wu, Xiang Li, Shangguang Wang, in *USENIX Annual Technical Conference (USENIX ATC, CCF-A)*, 2024.

[C6] "Mobile Foundation Model as Firmware"

Jinliang Yuan*, Chen Yang*, **Dongqi Cai***, Shihe Wang, Xin Yuan, Zeling Zhang, Xiang Li, Dingge Zhang, Hanzi Mei, Xianqing Jia, Shangguang Wang, Mengwei Xu, in *Proc. ACM Int. Conf. Mobile Computing and Networking (MobiCom, CCF-A)*, 2024.

[C5] "Federated Few-shot Learning for Mobile NLP"

Dongqi Cai, Shangguang Wang, Yaozong Wu, Felix Xiaozhu Lin, Mengwei Xu, in *Proc. ACM Int. Conf. Mobile Computing and Networking (MobiCom, CCF-A)*, 2023.

[C4] "Efficient Federated Learning for Modern NLP"

Dongqi Cai, Yaozong Wu, Shangguang Wang, Felix Xiaozhu Lin, Mengwei Xu, in *Proc. ACM Int. Conf. Mobile Computing and Networking (MobiCom, CCF-A)*, 2023.

[C3] "GPT4D: Automatic Cross-Version Linux Driver Upgrade Toolkit"

Borui Yang, Hongyu Li, **Dongqi Cai**, in the *8th EAI International Conference on Machine Learning and Intelligent Communications (MLICOM)*, 2023.

[C2] “FedAdapter: Efficient Federated Learning for Mobile NLP”

Dongqi Cai, Shangguang Wang, Yaozong Wu, Mengwei Xu, in *Proceedings of the ACM Turing Award Celebration Conference (TURC)*, 2023.

[C1] “Mitigating App Collusion using Machine Learning”

Xuefei Duan, Hua Lu, Jinliang Yuan, Qiyang Zhang, **Dongqi Cai**, in *IEEE 7th International Conference on Big Data Intelligence and Computing (DataCom)*, 2021.

Journal Publications (* = equal contributions)

[J3] “Resource-efficient Algorithms and Systems of Foundation Models: A Survey”

Mengwei Xu* (My advisor), **Dongqi Cai***, Wangsong Yin*, Shangguang Wang, Xin Jin, Xuanzhe Liu, accepted in *ACM Computing Surveys (ACM CSUR, Impact Factor: 23.8, ranked 1/143 in Computer Science Theory & Methods)*, 2024.

[J2] “Accelerating Vertical Federated Learning”

Dongqi Cai, Tao Fan, Yan Kang, Lixin Fan, Mengwei XU, Shangguang Wang, Qiang Yang, e in *IEEE Transactions on Big Data (IEEE TBD)*, 2024.

[J1] “Implementation of an E-payment security evaluation system based on quantum blind computing”

Dongqi Cai, Xi Chen, Yuhong Han, Xin Yi, Jinping Jia, Cong Cao, Ling Fan, in *International Journal of Theoretical Physics (IJTP)*, 2020.

Workshop Publications (* = equal contributions)

[W4] “Large Language Models on Mobile Devices: Measurements, Analysis, and Insights”

Xiang Li, Zhenyan Lu, **Dongqi Cai**, Xiao Ma, Mengwei Xu, in *Proceedings of the Workshop on Edge and Mobile Foundation Models (EdgeFM)*, co-located with *ACM International Conference on Mobile Systems, Applications, and Services (MobiSys, CCF-B)*, 2024.

[W3] “FedRDMA: Communication-Efficient Cross-Silo Federated LLM via Chunked RDMA Transmission”

Zeling Zhang*, **Dongqi Cai***, Yiran Zhang, Mengwei Xu, Shangguang Wang, Ao Zhou, in *Proceedings of the 4rd Workshop on Machine Learning and Systems (EuroMLSys)*, co-located with *European Conference on Computer Systems (EuroSys, CCF-A)*, 2024.

[W2] “Towards Practical Few-shot Federated NLP”

Dongqi Cai, Yaozong Wu, Haitao Yuan, Shangguang Wang, Felix Xiaozhu Lin, Mengwei Xu, in *Proceedings of the 3rd Workshop on Machine Learning and Systems (EuroMLSys)*, co-located with *European Conference on Computer Systems (EuroSys, CCF-A)*, 2023.

[W1] “Towards ubiquitous learning: A first measurement of on-device training performance”

Dongqi Cai, Qipeng Wang, Yuanqiang Liu, Yunxin Liu, Shangguang Wang, Mengwei Xu, in

*Proceedings of the 5th International Workshop on Embedded and Mobile Deep Learning (EMDL), co-located with ACM International Conference on Mobile Systems, Applications, and Services (**MobiSys, CCF-B**), 2021.*

Patents

[P4] "A Federated Learning Method, System, and Apparatus Based on Forward Gradient"

Mengwei Xu; Yaozong Wu; **Dongqi Cai**; Shangguang Wang

[P3] "A Federated Few-Shot Learning Method, System, and Device for Natural Language Models"

Mengwei Xu; **Dongqi Cai**; Ao Zhou; Xiao Ma; Shangguang Wang

[P2] "A Federated Learning Method, Device, and System for Pre-trained Models"

Mengwei Xu; **Dongqi Cai**; Ao Zhou; Xiao Ma; Shangguang Wang,

[P1] "Vertical Federated Learning Modeling Optimization Method, Device, Medium, and Program"

Dongqi Cai; Lixin Fan; Qiang Yang

Teaching Experience

- Teaching Assistant, Principles of Machine Learning Systems, University of Cambridge (Michaelmas Term 2024)

Invited Talk

- EMDL'21 (Co-located with MobiSys'21), Towards ubiquitous learning: A first measurement of on-device training performance, Online, 2021/06/25
- EuroMLSys'23 (Co-located with EuroSys'23), Towards Practical Few-shot Federated NLP Rome, Italy, 2023/05/08
- MobiCom'23, Efficient Federated Learning for Modern NLP, Madrid, Spain, 2023/10/05
- MobiCom'23, Federated Few-shot Learning for Mobile NLP, Madrid, Spain, 2023/10/05
- Northwestern Polytechnical University, PhD Research Methodology, Online, 2023/10/30
- BUPT 'Diligent Research, Academic Leadership' Academic Forum, Efficient Federated Learning for Modern NLP, Beijing, China, 2023/12/26
- EuroMLSys'24 (Co-located with EuroSys'24), FedRDMA: Communication-Efficient Cross-Silo Federated LLM via Chunked RDMA Transmission, Athens, Greece, 2024/04/22
- MoiSys'24 N2Women, Large Language Models on Mobile Devices: Measurements, Analysis, and Insights, Tokyo, Japan, 2024/06/03
- EdgeFM'24 (Co-located with MobiSys'24), Large Language Models on Mobile Devices: Measurements, Analysis, and Insights, Tokyo, Japan, 2024/06/07
- USENIX ATC'24, FwdLLM: Efficient Federated Finetuning of Large Language Models with Perturbed Inferences, SANTA CLARA, CA, USA, 2024/07/11
- AI TIME NeurIPS 2024 Forum, SILENCE: Protecting Privacy in Offloaded Speech Understanding on Resource-constrained Devices, Online, 2024/11/20
- NeurIPS'24, SILENCE: Protecting Privacy in Offloaded Speech Understanding on Resource-constrained Devices, Vancouver, Canada, 2024/12/11
- CCF Talk, Efficient Federated Learning System for LLMs, Online, 2024/12/22